

Cross-Lingual Transfer Learning for Low-Resource NLP

David Lee¹², Eva Novak²³, Frank Schmidt¹

¹Technical University of Munich, Germany

²Max Planck Institute for Intelligent Systems, Germany

³University of Edinburgh, UK

{dlee, fschmidt}@tum.de | e.novak@ed.ac.uk

Abstract

We present a multilingual pre-training approach that achieves strong transfer from high-resource to low-resource languages. Using a shared subword vocabulary of 128k tokens, we train on 50 languages simultaneously.

1 Introduction

Low-resource languages account for the majority of the world's 7000+ languages yet receive less than 1% of NLP research attention. Cross-lingual transfer offers a promising path: pre-train on data-rich languages and fine-tune on minimal target-language data.

Acknowledgements

David Lee was funded by DFG grant LE-2345/1-1.

Eva Novak acknowledges support from the MPI-IS International Max Planck Research School for Intelligent Systems (IMPRS-IS).

Compute was provided by the Leibniz Supercomputing Centre.

Contact: research-group@tum.de

References

[1] Conneau et al. Unsupervised cross-lingual representation learning. NeurIPS 2020.

[2] Pires et al. How multilingual is multilingual BERT? ACL 2019.